

A platform for Internet-connected wireless sensor networks

Stuart Whitehead

3rd January 2016

Abstract

Acknowledgements

Contents

Abstract	ii
Acknowledgements	iii
Contents	iv
List of Figures	vi
List of Tables	vii
1 Introduction	1
2 Background	2
2.1 Opportunities	3
2.2 Challenges	5
2.3 Use Cases	7
2.4 Existing Platforms	8
2.4.1 Overview	8
2.4.2 Hosting	9
2.4.3 Source Availability Model	9
2.4.4 Connectivity	10
2.4.5 Bidirectional Communication	12
2.4.6 Application Triggers	13
2.5 Sensor Hardware	14
2.5.1 Overview	14
2.5.2 Processor Architecture	14
2.5.3 Memory	15
2.5.4 Input/Output	15
2.5.5 Connectivity	16
2.6 Case Study: UK Smart Meters	18
2.6.1 Overview	18
2.6.2 Governance	18
2.6.3 Implementation	19
2.7 Conclusions	20
3 Specification	21
4 Design	22

5	Implementation	23
6	Results and Evaluation	24
7	Future Work	25
8	Conclusions	26
	Bibliography	27

List of Figures

2.1	The boundary between IoT and other trends in computing	3
-----	------------------------------------------------------------------	---

List of Tables

2.1	A summary of the current state of the Internet of Things	8
2.2	Characteristics of existing IoT platforms	13
2.3	Summary of hardware boards	17

Chapter 1

Introduction

Chapter 2

Background

The term ‘Internet of Things’ captures the essence of the vision well; the vision of a world where everyday, physical objects are gateways to web-based services. These so-called ‘smart objects’ comprise of sensors to perceive their environment or context, as well as a notion of interconnectivity with other objects or services. Connected objects or services may react to collected data to trigger actions, and these actions may be digital or physical. The Internet of Things could be summarised as data collection, aggregation and reaction in the physical domain.

The boundaries between IoT and other trends help to define its place in computing. For instance the research area of Wireless Sensor Networks (WSN) carries similarities in hardware requirements and challenges. In particular, WSN also comprise of connected sensors and actuators (Mottola and Picco, 2011). WSNs differ from IoT because of the scope of connectivity—they are limited to specific applications and area coverage, such as a production line or a single building. In comparison, IoT applications have a global scope which encourages co-operation between devices, networks and services. Wireless Sensor Networks can, however, form one aspect of an IoT application.

The research area of ‘Wearables’ also overlaps with IoT. Wearables are ‘smart’ devices designed to be worn or embedded within the body and combine sensors and some form of connectivity, typically integrating with a smartphone (Wei, 2014). Consumer Wearables products are already on the shelves, such as Fitbit—a wrist-worn personal health tracker. The Fitbit wristband connects to a smartphone with Bluetooth Low Energy (BLE) and this smartphone then provides global connectivity through its wireless connection. Users can sign-in to a web-based dashboard which will collate and organise personal data. In this scenario, Wearables are collecting, aggregating and reacting to data in the physical domain. Wearables could therefore be described as a subset of IoT (as shown by Figure 2.1) because they are an application in a specific domain.

IoT can also take advantage of research conducted for the ‘Big Data’ trend. Over the past three decades, scalable and distributed data management have been a focus for the database research community (Agrawal et al., 2011). This research originally focused on traditional enterprise applications, but it has expanded to include the Cloud Computing paradigm across two classes of systems: update-heavy applications and ad-hoc analytics & decision support. IoT is expected to generate large volumes of data and so by leveraging Big Data analytical processes, valuable information can be deduced. For example, Fitbit’s web-based dashboard could use Big Data analytics to analyse personal health statistics and to find data trends.

IoT also goes some way to realise the Ubiquitous Computing concept. Ubiquitous Computing (also known as pervasive computing) envisages a world where users are surrounded by connected

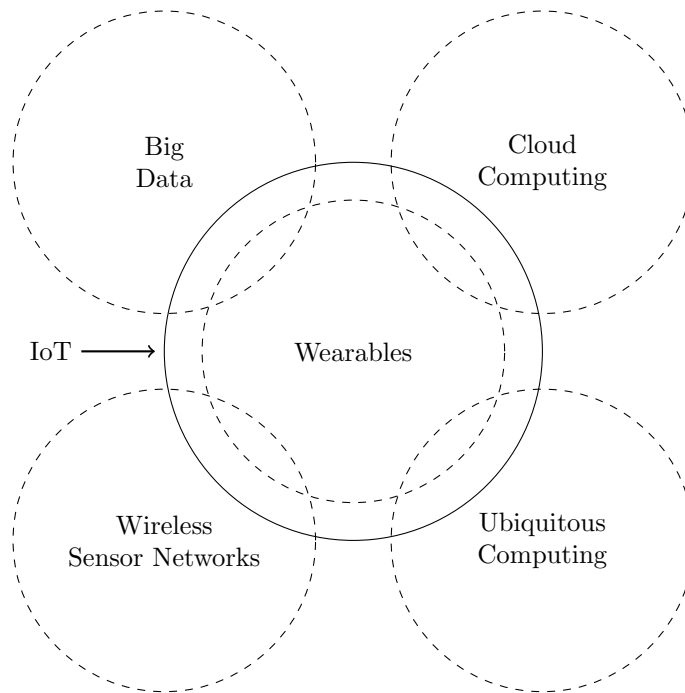


Figure 2.1: The boundary between IoT and other trends in computing

technology—technology in our homes, workplaces and recreational activities. In his seminal article, Weiser (1999) notes that “specialized elements of hardware and software, connected by wires, radio waves and infrared, will be so ubiquitous that no one will notice their presence”. Lyytinen and Yoo (2002) take this further and suggest that we will not only be surrounded by technology, but it will be embedded in our social and physical interactions. Research areas such as this not only contribute to the IoT concept but also benefit from its research endeavours.

2.1 Opportunities

Why should we develop Internet of Things solutions? Would our society actually benefit from them? These are the questions posed by businesses, consumers, software developers and academics alike. Pilot projects as well as industrial and academic research have shown that there are wide-reaching opportunities to be grasped. Until recently, these opportunities were achievable only in theoretical, lab-based scenarios but due to the advancements in the supporting technology, they are more achievable for real-world industries.

The technical capabilities of hardware, the readiness of software and all of their associated costs are notable impediments to the IoT industry. Advancements in the hardware—smaller, more powerful and more efficient microcontrollers—mean that smart objects are technologically more feasible. Large industry players, such as Intel and Samsung, have brought to market the their own IoT hardware platforms—Intel’s Edison System on a Chip (SoC) and Samsung’s Artik family of boards. Improvements in the mass production of these also reduce financial barriers, making investment more feasible for businesses. The real-world opportunities presented before us can be categorised into two main areas: economic & societal and technical.

Economic and Societal

Irrespective of the industry or domain, the ultimate purpose of Internet of Things is to help people. From the perspective of a service, these applications do and should provide value on an individual basis, however the real value is found with the network effect of interconnected products and services. Both the Internet of Things and societies of people are similar in that their wholes are greater than the sum of their parts.

The European Union (EU) is a prime example of a society which can benefit from the Internet of Things. It is special because it represents the needs and wants of not just one nation but a collection of nations. The European Commission is the governmental body of the EU, which itself recognises the potential in IoT: “One major next step in this development [of the Internet] is to progressively evolve from a network of interconnected computers to a network of interconnected objects, from books to cars, from electrical appliances to food, and thus create an ‘Internet of things’.” In 2009 the Commission published an action plan for embracing the Internet of Things.

The action plan notes two main areas of opportunity: citizen well-being and economic prosperity. The improved well-being of citizens is achieved through specific and targeted use-cases. For instance Internet-connected health monitoring systems could alleviate pressure on medical services for the ageing society, or smart waste management with products which can describe their contents would help to reduce their carbon footprint. The improvement of economic prosperity is achieved through organised and systematic uptake of the Internet of Things. By leading the development of IoT rather than accepting the standards of other nations, the EU can drive development for the benefit of its own industries.

Although there are strong opportunities for individuals and political territories, the best economic opportunities are available to businesses. The Internet of Things has the potential to provide many new income streams to businesses as well as finding other money through efficiency savings. First of all IoT opens up new markets which were previously not feasible nor even considered. These markets may include consumer products such as the previously mentioned Fitbit or novel service solutions such as home security and monitoring. For organisations with complex supply chains or processes, the Internet of Things could help to maximise resources and stock control. This is especially evident in areas such as logistics or public transport where connected devices could help to reduce costs by optimising transportation routes automatically.

Technical

Technical opportunities and advancements are self-perpetuating and really are the driving force behind the Internet of Things. As previously noted, IoT would not be realistically possible without improvements in the underlying technology. As more powerful, efficient and cheaper hardware boards are delivered and as they become supported by software layers, the possible use-cases of them diversifies too.

Mattern and Floerkemeier (2010) note that the Internet of Things is not the result of a single novel technology but rather several complementary technical developments. These technical developments provide a range of capabilities, or in the eyes of an application stakeholder, technical *opportunities*. The main opportunities presented are: communication and cooperation; addressability; identification; sensing; actuation; localisation; embedded information processing and user interfaces. The authors also note that most applications will need only a subset of these capabilities, but it is better to have the option there.

The value of the Internet of Things is derived from the communication of data—either human-to-machine, machine-to-human or machine-to-machine. This trait makes the capabilities of communication and cooperation, addressability and identification particularly important to its success. Objects with the ability to network between themselves or other Internet services

are clearly key. While the Internet will provide the backbone for this, it is the Wireless Personal Area Network technologies which present the opportunities—technologies such as GSM, Wi-Fi, Bluetooth, ZigBee and 6LoWPAN. These technologies in conjunction with other technology layers allow devices to be uniquely identified and addressed from anywhere in the world. With this global interconnectivity, the possibilities really do open up.

The second main characteristic of IoT is the ability to interact with the physical domain. Digital applications will interact with the physical world in two ways: through the monitoring and sensing of physical properties and through the actuation of the physical world in reaction to data. There are a massive variety of sensors available, even to hobby markets, and can measure any imaginable physical property—such temperature (ambient or spot), gas particulates, pressure and distance. Of course sensor hardware is nothing new, however the improved support and reduced cost opens up further opportunities. In a similar vein, actuators such as motors, solenoids or lamps are nothing new but when combined with IoT boards in consumer appliances or industrial applications, anything is possible.

Advancements in technology even provide new opportunities in system architecture. Smart objects will generate a lot of data and given the projected increase in their numbers, the bandwidth available across Internet backbone would become wasted. Since IoT microcontrollers are becoming more powerful, there are opportunities for embedded information processing. Embedded information processing refers to these end devices performing some form of data processing or storage before transmitting their data; for example, an end device might analyse its own data and only transmit if a specific threshold is met, rather than relying on cloud computing-based processing. This is an active area of research called edge computing (or fog computing) and allows for novel methods of data processing.

2.2 Challenges

The previously mentioned papers provide solid arguments for investing in the Internet of Things, however this is still an active area of research with numerous challenges to be addressed. Before IoT will become a mainstream paradigm, the technologies and ecosystem surrounding them must settle and become more mature. In a similar way to opportunities, the challenges can be categorised as Economic & Societal or Technical.

Economic and Societal

If the ultimate purpose of the Internet of Things is to help people then it would be meaningless without them. This is the risk that the industry is carefully attempting to balance; the social acceptance of IoT products is imperative to their success but consumers and markets could freak out if industries try too much too soon. The European Commission's action plan describes practical challenges which need to be addressed to develop this acceptance: governance, standardisation, security, data privacy and trust.

Many aspects of digital technology is governed by public bodies and the Commission believes that IoT should not be any different. For instance the Internet Assigned Numbers Authority (IANA) is responsible for global IP addressing and the DNS root, two critical components of the Internet. The Commission notes that technology will advance regardless of public intervention due to a normal cycle of innovation and that “simply leaving the development of IoT to the private section... is not a sensible option in view of the deep societal changes that IoT will bring about.”

The main areas which may need to be governed are identification, information security and ethical & legal accountability. As we have already established, the value of the Internet of

Things is driven by the interconnectivity between devices. Various mechanisms already exist to uniquely address IoT devices however these are not all compatible between applications; if they were, should a public body be responsible for assigning unique identifiers, similar to the Media Access Control (MAC) addresses assigned by the Institute of Electrical and Electronics Engineers (IEEE)? What if an entity illegally or immorally handled sensor or user data—how can they be held accountable and for what? The Internet of Things as a whole will be affected by ineffective governance of these areas—in particular it will suffer from stifled innovation; a mistrust with data and jeopardised system integrity.

Standardisation is a technical consideration but also impacts on the economic potential of IoT. The purpose of a standard is to give different entities and stakeholders a common language to design, build or communicate. The widespread adoption of a standard gives businesses of all sizes the opportunity to develop interoperable solutions. In being interoperable, standards-based solutions will be suitable for a wider market and in turn, this will encourage innovation and improve international competitiveness. If the Internet of Things is to maximise economic impact and to enjoy widespread adoption, there must be accepted practices or formal standards in place.

While governance and standardisation will support IoT, users must be able to accept this new paradigm in their own time. The greatest challenges in this respect are security and data privacy—the protection of privacy and personal data are two fundamental rights in the European Union. The challenge with this, in respect to IoT, is that new methods of collecting and using data will be invented. Does current legislation and safeguards protect the interest of EU citizens adequately? Who will own the data—the user whom it is about or the organisation that captured it? The Commission suggests that in response to this challenge, IoT components should be designed from their inception with a privacy- and security-by-design mindset.

Technical

These challenges can be cross-dependant and many of the points raised about economic & societal challenges are dependant on technical issues being addressed. Since the Internet of Things is not a single novel technology but a collection of cooperating tools, the range of technical challenges is diverse. They vary from the aspects of user experience; device interoperability and discovery; system complexity with regards to scalability, data management & interpretation and code-level complexity as well as hardware challenges covering fault tolerance, power supply and wireless communication.

A good user experience is important for the adoption of IoT products however the underlying technology currently makes some aspects of this difficult. Mattern and Floerkemeier (2010) note that since smart objects will be used sporadically, they “need to establish connections spontaneously, and organize and configure themselves to suit their particular environment.” They coin this requirement ‘Arrive and Operate.’ A typical home network might use a wireless router, such as the BT HomeHub, to provide wireless local area network (WLAN) coverage. To connect to this network, users are expected to enter some form of Pre-Shared Key (PSK) on the device. While this user experience is fine for computers, laptops and smartphones it is less than ideal for smart objects with little or no user interface. The challenge here is to develop technology which allows consumers to use smart objects with little or no configuration.

IoT innovation is expected to deliver a diverse range products and solutions which makes application interoperability and discovery a challenge. Since manufacturers have the freedom to implement proprietary tools, two systems developed by different manufacturers may not be immediately compatible and this reduces the overall potential of the concept. This issue is further exacerbated by the variety in hardware processing and communication capabilities. The Internet of Things therefore requires common practices and standards to be accepted by the industry (as

previously mentioned with regards to economic side-effects). Research efforts have gone some way to address these challenges, such as the IoT ‘meta-hub’ by Mineraud and Tarkoma (2015).

By all estimates, the Internet of Things will have a larger scope and deployment footprint when compared to the existing Internet of computers which needs to be addressed. The European Commission puts this figure at 50–70 billion devices (on average, 10 per human). Mattern and Floerkemeier (2010) also note that things will be cooperating mainly within a local environment. This means that IoT devices, software and infrastructure must work equally efficiently in both small- and large-scale environments. This complexity is also reflected with the volume of unstructured data being generated; data analysis must be able to scale also, as is being addressed with research around Big Data.

On a lower level, the main challenges present are fault tolerance, power supply and wireless communication. Given how variable IoT environments are (office spaces, homes, public spaces, remote and exposed areas) and how complex the supporting architectures could be, resilience to faults is important. IoT applications should therefore have redundancy across its various technical layers. One main cause of technical faults is the power supply driving any given IoT object. ‘Things’ are typically mobile and therefore need a self-sufficient power source; the industry is therefore challenged to produce long-lasting batteries while maximising power efficiency in other areas. Existing wireless communication technologies are too bloated and consume too much power, such as GSM, Wi-Fi and Bluetooth. To tie these challenges of fault tolerance, power consumption and wireless communication together, IoT applications need lightweight and robust wireless communication standards.

2.3 Use Cases

Use-cases are a helpful mechanism for demonstrating the Internet of Things concept. It is a difficult concept to explain because there are many cooperating components and ideas and for this reason, research papers exemplify their arguments with practical examples. Although IoT applications could be developed for virtually any industry there are some areas where its application is more obvious. Some areas which are repeatedly brought up by researchers are smart cities, utilities and logistics. Assuming that the challenges (summarised in Table 2.1) are met, these areas are very achievable.

‘Smart Cities’ is a vision which convincingly demonstrates many benefits of IoT. The Department for Business Innovation & Skills (2013) describes the main issues facing local authorities, including: piecemeal urban infrastructure, climate change targets, the changing nature of the high street and elderly social care. The economic downturn has also reduced budgets assigned to local authorities by as much as 30% between 2010 and 2013 making cost efficiencies another issue. By developing or retrofitting urban infrastructure with IoT connectivity, local authorities can monitor services in much finer detail. For instance, Bigbelly Solar is a smart waste and recycling system. The Internet-connected waste bins allow local authorities to monitor their status and to schedule collections when they are full. Bonomi et al. (2012) also describe a system of smart, connected traffic lights which can control the traffic flow through a whole city, reducing congestion and improving safety.

IoT in the utilities sector is one use-case which has started to become realised. The British Government is requiring energy companies to install smart meters for their customers and expect most to have a smart meter installed by 2020 (The Department of Energy & Climate Change, 2013). Smart meters monitor domestic energy usage on a per-site or per-socket basis and can give real-time information how much electricity is being used, expressed in pounds and pence. For individual households this has the benefit of enabling homeowners to save money and to reduce

Overlapping Research	Opportunities		Challenges		Example Use Cases
	Economic & Societal	Technical	Economic & Societal	Technical	
Cloud Computing	Citizen Well-Being	Communication & Cooperation	Governance	‘Arrive and Operate’	Smart Cities
Big Data	Economic Prosperity	Addressability	Standardisation	Interoperability	Waste Management
Ubiquitous Computing	Ageing Society	Identification	Security	Discovery	Smart Traffic Lights
Wireless Sensor Networks	New Markets	Sensing	Data Privacy	Software Complexity	Utilities
Wearables	Business Optimisation	Actuation	Trust	Scalability	Smart Meters
Edge Computing		Localisation		Data Management	Logistics
		Embedded Information Processing		Fault Tolerance	
		User Interfaces		Power Supply	
				Wireless Communication	

Table 2.1: A summary of the current state of the Internet of Things

emissions. At a national level, smart meters are a step towards reducing energy consumption and meeting climate change goals. Smart thermostats like Nest also allow households to minimise their gas consumption and to reduce heating bills.

Logistics is one area which demonstrates the efficiency savings that an IoT application can make. The European Commission (2007) states that “Freight Transport Logistics focuses on the planning, organisation, management, control and execution of freight transport operations in the supply chain” and notes that it is one of the drivers of European competitiveness. To remain competitive, this industry must continue to improve and streamline infrastructure, fleet management and goods tracking. The Internet of Things could be used to automate good identification and location with inexpensive RFID tags and GPS chips as well as optimise transport routes. This would reduce time spent by workers manually recording goods and it will reduce errors at interchange points. The Commission has coined this the ‘Internet for cargo’.

2.4 Existing Platforms

2.4.1 Overview

Now that a baseline has been established for Internet of Things research, it is possible to investigate specific technologies and architectural decisions. There are already numerous IoT platforms both in service and under active development; in fact, Amazon announced their own platform (Amazon IoT) when writing this very comparison and not long after, Digi refreshed their own platform product, showing just how dynamic this area currently is.

The aim of an IoT platform and the tools which it provides does differ from vendor to vendor but after comparing some headline features, there are various characteristics common amongst them. The platforms were chosen by sifting through marketing websites online and paper lists, like that by Mineraud et al. (2015), before selecting a representative sample of organisations. The platform offerings were analysed and compared against five common characteristics: hosting model, source availability model, connectivity, bidirectional communication support and trigger support.

2.4.2 Hosting

The hosting model offered by an organisation defines which stakeholder is responsible for maintaining the IoT application and its supporting infrastructure—the platform provider or the customer. This is notable differentiator between platforms and can be indicative of the platform vendor’s business objectives (which also relates to the Source Availability Model, coming up next). There are generally two variations in hosting models: Platform as a Service (PaaS) or self-hosting.

Platform as a Service

A Platform as a Service (PaaS) is a Cloud Computing concept. It is an abstraction of the complete application technology stack including both hardware and software. An organisation providing a PaaS would be responsible for the management and maintenance of the hardware such as servers, RAID backing storage and networking appliances like routers or switches. Depending on the software on offer, the vendor would also be responsible for managing and maintaining the server operating system as well as any supporting software packages and applications. Customers are typically provided access to this abstraction through a control panel or an Application Program Interface (API).

In the context of the Internet of Things, a PaaS would provide various benefits to a customer. Primarily the customer do not need to allocate resources for infrastructure maintenance since the vendor manages this, with the result being more time for the customer to focus on business needs. There is also less investment required on behalf of the customer—PaaS resources can be provisioned when and if they are needed, rather than purchasing hardware upfront. Another benefit is the extra layer of support—system scalability, stability and fault tolerance are the responsibility of the PaaS vendor. One example of a PaaS is Carriots.

Self-hosted

An alternative to a Platform as a Service is the self-hosted option. As the name implies, this option makes customers responsible for the maintenance and management the application infrastructure (although software updates may still be provided by the platform vendor). There are technical and legal arguments to support such a business decision.

The main technical aspect is that of control. With the self-hosted option, the system architecture can be optimised for specific use-cases, such as the global location of servers to reduce application latencies (the time taken for network transmissions between devices and server). This also gives customers direct access to the core software and makes it easier to modify it for their own purposes (if the licence or software allows for that).

Self-hosting may also be necessary for legal compliance, especially in the realm of data protection. Depending on the country of operation, it may be necessary to retain data within country boundaries and this may not be enforceable with a PaaS. For example, in October 2015 the Court of Justice of the European Union ruled that the Safe Harbour agreement between the EU and United States is invalid (Case C-362/14). This means that any organisation sharing data under this agreement is now acting unlawfully—a problem which could be avoided if data is managed from within the European Union. One example of a self-hosted platform is Nitrogen.io.

2.4.3 Source Availability Model

The Source Availability Model chosen by a platform developer defines the level of accessibility for their source code. This model comes in two forms—open or closed source—and can be further

refined by a software license framework.

The source code of Open Source Software (OSS) is made publicly available. This could be served as a standard file download or, as is becoming common practice, on a code sharing service such as Github or Bitbucket. OSS will typically be provided under the conditions defined in a license, such as the GPL, MIT or Apache license frameworks. These vary on various points like distribution, modification and notification of copyright (Github, Inc., 2015).

Closed Source Software is proprietary source code not made publicly available, although it may be necessary to distribute compiled binaries. This just means that the organisation does not wish to make source code available for modification or knowledge sharing.

In conjunction with the hosting model, the source availability is indicative of business objectives. Organisations wishing to commercialise their platform as a product offering will use a Closed Source model to protect their intellectual property. Conversely those organisations where knowledge sharing is an aim, or where the platform is a mechanism to sell other products or services (such as Sparkfun and their electronic hardware) may use an Open Source model.

2.4.4 Connectivity

The Internet of Things concept *is* interconnectivity and communication, making device and application connectivity a very important topic. The platforms investigated support a range of connectivity protocols and concepts. These range from widely-accepted HTTP-based APIs to more niche protocols like CoAP and MQTT.

HTTP

The Hypertext Transfer Protocol (HTTP) underpins the internet. It benefits from widespread support and most notably, is used by web browsers to interface with web servers. RFC 2616 (Fielding et al., 1999) defines the specification for HTTP version 1.1 and notes:

It is a generic, stateless, protocol which can be used for many tasks beyond its use for hypertext, such as name servers and distributed object management systems, through the extension of its request methods, error codes and headers.

As hinted at by this RFC 2616 quote, three important features of HTTP are request methods, error codes and headers. To generalise the HTTP life cycle, a user agent will send a request to an origin server, the server will process the request and will return a response. The request will use an HTTP *verb* to define the action to take, such as `GET` or `POST`, and will include headers to specify other parameters. The response has an associated status code, such as `200 OK` or `404 Not Found`.

Most of the platforms investigated support HTTP requests through a Representational State Transfer (REST) API. Fielding (2000) states that the REST architectural style emphasises scalability, independence of components, minimises latency and maximises security. These characteristics are achieved by applying constraints to the application, such as stateless communication, a uniform component interface and a layered approach to infrastructure technologies. REST methods are based on the HTTP verbs, such as `GET`, `POST`, `PUT` and `DELETE`. For IoT application developers, a REST API provides a scalable, uniform and robust interface between devices and the Internet.

Software engineers working on specifications for the Internet have foreseen the Internet of Things, even if it was taken in jest. RFC 2324 (Masinter, 1998) is an April Fool's joke and defines the Hyper Text Coffee Pot Control Protocol version 1.0 (HTCPCP/1.0). It specifies methods by which an Internet-connected coffee pot can be controlled, and even adds a new error

code: “Any attempt to brew coffee with a teapot should result in the error code ‘418 I’m a teapot’. The resulting entity body MAY be short and stout.”

WebSockets

The WebSocket Protocol is an independent TCP-based protocol that enables two-way communication. It uses an HTTP request as part of the handshaking process, which is treated as a connection upgrade request. RFC 6455 (Fette and Melnikov, 2011) states “The goal of this technology is to provide a mechanism for browser-based applications that need two-way communication with servers that does not rely on opening multiple HTTP connections (e.g., using `XMLHttpRequest` or `<iframe>`s and long polling).” This technology plays a vital role in enabling bidirectional communication between web browsers and servers, as will be covered in section 2.4.5.

WebSockets use standard HTTP ports for communication by default (port 80 and 443 for unencrypted and encrypted connections respectively). This has the added benefit of being compatible with common network and firewall configurations.

MQTT

MQ Telemetry Transport (MQTT) is a simple and lightweight messaging protocol. It is based on a publish/subscribe model and is specifically designed for constrained devices and low-bandwidth, high-latency or unreliable networks. In general, MQTT is used in conjunction with a message broker. Clients subscribe to topics on this broker which will then forward any relevant messages. Clients are then free to do whatever they please with these messages (Banks and Gupta, 2014).

MQTT uses TCP/IP to transmit messages and operates on ports 1883 and 8883 for unsecured and secured uses respectively. Since these ports are not common, firewalls will typically block access. For this reason, MQTT can be tunnelled using WebSockets (and means this can be used in the browser, too).

CoAP

Constrained Application Protocol (CoAP) is another protocol designed with the Internet of Things in mind. Shelby et al. (2014) describe CoAP as “a specialised web transfer protocol for use with constrained nodes and constrained networks.” Some headline features align with principles used with HTTP, such as utilisation of the REST model. RFC 7252 defines the proposed specification for CoAP. It specifically notes some main features:

- Web protocol fulfilling M2M requirements in constrained environments
- UDP [RFC0768] binding with optional reliability supporting unicast and multicast requests
- Asynchronous message exchanges
- Low header overhead and parsing complexity
- URI and Content-type support
- Simple proxy and caching capabilities

CoAP is likened to HTTP throughout the specification however there are some notable differences. HTTP uses the TCP/IP transport protocol whereas CoAP uses UDP, a protocol with significantly less overhead. UDP is not a request/response protocol so to provide a request/response interaction model, CoAP must deal with this through a logic layer.

UDP

One of the researched platforms also offers the User Datagram Protocol (UDP) as a connectivity option. RFC 768 (Postel, 1980) notes “this protocol provides a procedure for application programs to send messages to other programs with a minimum of protocol mechanism.” To achieve ‘a minimum of protocol mechanism’, UDP does not establish a connection with the receiving device, so unlike TCP/IP it is connectionless. The protocol also specifies no mechanism for error correction so if a packet is lost in transit, it is lost forever.

Since data can be lost, this protocol should only be used in applications which can tolerate data loss. To give a real-world example, a battery-powered hygrometer sensor may be collecting data about the environmental humidity. If UDP was used as the connectivity protocol, it would allow this sensor to wake up, transmit its reading and then fall back to sleep without being concerned about the packet’s progress. If a packet was lost, it would not have business or safety impacts.

2.4.5 Bidirectional Communication

One main opportunity of IoT, as described in Section 2.1, is to drive actuators in response to data. There are three aspects which define whether unidirectional or bidirectional communication is possible: the connectivity protocol, the software application and the network environment. To illustrate the following points, a hypothetical scenario will be used, where one server and one actuator are connected together through the internet. In this case, unidirectional communication means that only the actuator can initiate network requests whereas bidirectional communication allows either the server or actuator to initiate.

If the actuator is publicly addressable with a dedicated IP address, bidirectional communication is trivial to implement. Supposing that the HTTP protocol is being used for communication, both the actuator and server can initiate requests to each other by using the appropriate IP address. The server may send an HTTP request instructing the actuator to do something, and it obey. In real-world scenarios, there are three factors which make bidirectional communication like this difficult: firewalls, Network Address Translation and IPv6 uptake.

Firewalls are common-place both in businesses and at home. They are a network security tool and can be either a dedicated hardware unit or software built into other devices, such as a domestic router. Their purpose is to screen both in and outbound connections which may have a malicious or compromising effect on the network integrity. The side effect which this has on IoT systems is that non-common protocols may be blocked, such as MQTT (since it uses ports 1883 and 8883). While devices behind a firewall may be publicly addressable, they may not be reachable, depending on the IoT application.

While obtaining a public IP address for each IoT device is a worthy aim, it is not achievable in many situations due to Network Address Translation. Network Address Translation (NAT) is a tool used to delay the depletion of IPv4 addresses and is very common with domestic internet connections. RFC 2663 (Srisuresh and Holdrege, 1999) states “NAT devices are used to connect an isolated address realm with private unregistered addresses to an external realm with globally unique registered addresses.” In essence, this allows a number of privately addressed devices to communicate through one public address. In terms of IoT, this means that any devices connected behind NAT will not be publicly addressable—the server could not directly send a request to the actuator.

Internet Protocol version 6 (IPv6) will solve the issue of IP address depletion. It will offer over 340 undecillion possible addresses, compared to just under 4 billion available from IPv4, and will allow every connected device to have its own IP address for the foreseeable future (The European Commission, 2008). With regards to the Internet of Things, this is still not achievable.

Service	Hosting	Source	Connectivity	Bidirectional	Triggers
Amazon IoT	PaaS	Closed	MQTT, HTTP	Yes	Yes
Bug Labs Dweet	PaaS	Closed	HTTP	Yes (app)	No
Carriots	PaaS	Closed	HTTP, MQTT	No	Yes
Sparkfun	PaaS or self-hosted	Open	HTTP	No	No
Exosite	PaaS	Closed	CoAP, HTTP, UDP	No	Yes
Google Brillo	PaaS	Closed	Unknown	Unknown	Unknown
Digi	PaaS	Closed	HTTP	Yes	Yes
IFTTT	PaaS	Closed	HTTP	No	Yes
Newaer	PaaS	Closed	HTTP	No	Yes
Nimbits	Self-hosted	Open	HTTP	No	Yes
nitrogen.io	Self-hosted	Open	MQTT, HTTP	Yes	Yes
ThingSpeak	PaaS or self-hosted	Open	HTTP	No	Yes

Table 2.2: Characteristics of existing IoT platforms

IPv6 is still not supported by many internet service providers (ISPs) and networks, including that of Robert Gordon University. For the Internet of Things, this means that universally supported IPv6 is still a while off.

There are some useable techniques to overcome these three issues and to enable (or mimic) bidirectional communication. The first of these solutions is HTTP polling. With standard polling, the actuator would periodically ask the server for updates. If this period was 1 second, the communication may feel like it was real-time and bidirectional. This is however wasteful because not every request will result in an update. A more efficient form of polling is called long polling. The actuator will send a request to the server, but if there are no updates yet, the server will not respond. When the server does have a message to send to the actuator, it will respond and then end the request. This is better, but implementations suffer from the added complexity of managing connections, and it is using the HTTP specification in a way it was not designed.

A final and more effective solution is to implement communication using WebSockets. As previously described, the WebSocket protocol allows for bidirectional, full duplex communication and it can solve every challenge already addressed. WebSockets use the same communication ports as HTTP (port 80 and 443) by default, meaning that firewalls will allow them. Networks that support web browsing will support WebSockets. Also, the WebSocket connection can be initiated by the client and still allow bidirectional communication. In this case, the actuator can open a WebSocket connection to the server and overcome any Network Address Translation in place because it does not require a public IP address.

2.4.6 Application Triggers

One application-level feature common across investigated platforms is some form of event trigger. Where supported, platforms can trigger actions based on the data received from sensors and third-party APIs. Combining examples already mentioned, a hygrometer may send data to an IoT application. An application trigger could be configured to turn on a dehumidifier actuator if the humidity passes a certain threshold. If bidirectional communication is supported, this trigger would react in real-time. Triggers could also be configured to interact with other third-party software APIs.

2.5 Sensor Hardware

2.5.1 Overview

The Internet of Things is a multidisciplinary area of research and is just as dependant on advancements in electrical and electronics engineering as it is computer science. Since the hardware driving IoT sensor nodes is intertwined with the upper layers of software, it would be appropriate to investigate the current state of this area too. Rather than research the intricacies of how these pieces of hardware work, this paper will investigate existing end device hardware with regards to their implementation in IoT applications, particularly technical capability and energy efficiency.

With a similar method to Section 2.4, various makes and models of hardware boards were analysed to find comparable characteristics. The manufacturers chosen—Arduino, Raspberry Pi, Intel and Samsung—represent a cross-section of the field. These manufacturers aim their products at different markets, from the DIY ‘maker’ market of the Arduino to the professional, consumer electronics market of Samsung. This section will compare and contrast hardware models produced by these manufacturers across a range of key areas: processor architecture, memory, input/output options and wireless connectivity. Table 2.3 summarises each model and their specifications.

2.5.2 Processor Architecture

Just like any other computer system, the processor is the heart of the hardware board. It is responsible for running the operating system, orchestrating the various system components and for running the user application. The type and specification of the processor defines what can be processed and how quickly, which is ultimately reflected in the responsiveness and capability of the IoT device. Processor throughput is not the only consideration when scoping an IoT device however—energy consumption is also an important factor. The specification of the processor can be categorised into smaller areas: processor type, instruction set, number of bits, clock speed and the number of cores.

There are two broad types of processor systems, namely microcontrollers and microprocessors. A microcontroller (also known as an embedded controller) is a complete computer in one chip. A single microcontroller chip will typically combine a central processing unit (CPU), small amounts of program memory and Random Access Memory (RAM) and input/output lines (Ibrahim, 2011). In comparison, a microprocessor requires other system components to function, such as external memory, input/output lines and a clock circuit. In this regard, microprocessor hardware cannot be upgraded because it is integrated whereas a microprocessor system can, but they are cheaper to manufacture.

Processor architecture can also vary between makes and models. Ibrahim notes that processors can be classified by the number of bits they process, with 8 bits being the most popular for microcontrollers. Processors with a greater number of bits are more powerful but are also more expensive. The instruction set used also effects the performance of the processor; Reduced Instruction Set Computer (RISC) instructions take only one CPU cycle to complete whereas Complex Instruction Set Computer (CISC) instructions can take multiple cycles.

Clock speed and the number of cores are two final factors of a processor’s performance. The clock speed describes how many cycles the CPU executes in one second and is regulated by a crystal timer or internal digital clock oscillator (DCO) A greater clock frequency means more instructions which can be processed in a time period, and the more powerful the processor. Microcontrollers typically have clock speeds in the tens of Megahertz whereas microprocessors might run at hundreds of Megahertz or a number of Gigahertz. However due to physical limitations of hardware components, there is a practical limit in possible clock speed. In this case

more cores are added to the processor allowing it to evaluate multiple instructions concurrently. With regard to power consumption, Mikhaylov and Tervonen (2010) state that clock speed is one factor affecting the lifetime of a battery powered device so the CPU should be carefully considered for its application.

2.5.3 Memory

In any computer system, different types of memory are used to store program instructions and the temporary data associated with them. While these types of memory can be split into more discrete categories, there are generally two types of memory: non-volatile and volatile. Non-volatile memory, such as Read-Only Memory (ROM) or Programmable Read-Only Memory (PROM), is used to store program instructions and fixed user data. Non-volatile memory is used to store this kind of data because it is not lost when power is disconnected. In comparison, volatile memory loses its contents when power is removed and so it is used for temporary or intermediate data.

Memory characteristics also contribute to the capabilities and performance of the hardware system. Primarily, the available memory size defines how much data can be stored. The Arduino boards researched have 32KB of program memory available compared to the Gigabytes available on some Samsung Artik boards. This constrains the number of program instructions which can be stored and means the software running on the Samsung Artiks could be much more complex. Memory word length also has implications. The word length is the number of bits which a single unit in memory can store, and is commonly found to be 8, 16, 32 or 64 bits. Larger word lengths can store larger pieces of instructions or data, such as larger numbers.

2.5.4 Input/Output

The Input/Output (I/O) lines of a hardware board are the interface between sensors or actuators and the processor (or beyond). Without this physical connectivity, developers would not be able to capture information about the physical world, stopping the IoT concept in its tracks. This aspect of hardware boards is therefore important to the flexibility and future potential of IoT applications. To ensure compatibility between hardware platforms and external hardware like sensors, various industry standards have been developed.

Before any I/O standards are discussed though, it is first important to establish the importance of the operating voltage for an individual electronics circuit. The hardware boards analysed operate at various logic voltages; for instance most Arduino boards run at 5V whereas the native voltage of the Intel Edison is 1.8V. The standard logic voltage also effects I/O operation. Digital sensors may have their own, incompatible standard logic voltage—a microcontroller running at 5V would damage a sensor running at 3.3V. If directly interfacing between processors and sensors, a logic level shifter might have to be implemented so careful planning is required. Logic voltage is another aspect which affects the lifetime of a battery-powered node (Mikhaylov and Tervonen, 2010).

General Purpose Input Output (GPIO) pins are a blank canvas for engineers and developers. They are digital I/O pins with no default use and as their name suggests, can be configured as input or output. Since they are digital, GPIO pins can only read or write with 2 discrete values, logic high (1) or logic low (0), and these values usually respect the standard system voltage. Depending on the device, a number of GPIO pins might support Pulse Width Modulation (PWM). PWM is a technique for emulating analogue levels by pulsing the digital value at differing frequencies. Using a light bulb as an example, PWM could emulate the effect of dimming the light, like what could be achieved with an analogue potentiometer.

While GPIO pins allow for basic I/O operations, there is a need for communicating with more complex data. There are various available standards for interfacing between low-level electronics components. The datasheets for the selected model note some common standards, namely: Serial Peripheral Interface (SPI), Universal Asynchronous Receiver/Transmitter (UART), Inter-Integrated Circuit (I²C) and Integrated Interchip Sound (I²S). There are a variety of uses for these standards so this paper will go no further than noting that they exist.

2.5.5 Connectivity

Interconnectivity is yet one more important aspect to hardware designed for the Internet of Things. Since IoT devices are becoming increasingly mobile and wire-free, connectivity is being discussed almost exclusively in regards to wireless communication protocols and standards. As was discussed in Section 2.1, some IoT-specific wireless standards have already been established. Many of these technologies are supported by the various hardware platforms investigated, validating their purpose. The standards supported by these hardware platforms are IEEE 802.11 (Wi-Fi), Bluetooth and IEEE 802.15.4 (Low-Rate Wireless Personal Area Network).

Wi-Fi is a technology which has become synonymous with wireless connectivity and the Internet. The IEEE 802.11 standard states that the purpose of Wi-Fi is “to provide wireless connectivity for fixed, portable, and moving stations within a local area” where a local area could refer to a home, classroom or office space. Wi-Fi benefits from wide utilisation in all areas and industries so wireless coverage is likely to be in place already. Aust et al. (2012) state that the radio frequencies used by Wi-Fi (2.4GHz and 5GHz bands) are becoming crowded and are increasingly resulting in radio interference. This is being addressed in a new standard, IEEE 802.11ah, which makes use of the 900MHz frequency band. The two main potential advantages of this standard are greater range and less power consumption, which is beneficial to IoT devices. Even with this standard, however, end devices are impeded by the large overhead of upper layer protocols used in conjunction with Wi-Fi.

The Bluetooth standard has been repurposed as an effective wireless communication protocol for IoT devices. Prior to the introduction of Bluetooth 4.0 in 2010, it was primarily used in multimedia applications such as streaming stereo audio (Chang, 2014). With the introduction of Bluetooth 4.0 came new features, namely Bluetooth Low Energy (BLE), which consumes such a small amount of power that a sensor can run from a coin battery for months or years. For an average usage scenario, BLE could be expected to achieve a practical maximum range of 77 meters (Labs, 2015) which makes it ideal for use-cases such as body sensors and intelligent transport systems.

One final wireless connectivity standard supported by the investigated hardware is IEEE 802.15.4 Low-Rate Wireless Personal Area Networks (WPAN). The standard states that it “provides for ultra low complexity, ultra low cost, ultra low power consumption, and low data rate wireless connectivity among inexpensive devices” and covers the physical layer (PHY) and Medium Access Control (MAC) sublayer specifications. This IEEE standard has manifested into various competing technologies, in particular ZigBee and 6LoWPAN. ZigBee is a set of layers built on top of 802.15.4 and adds three important features: routing, ad-hoc network creation and self-healing mesh (Faludi, 2010). In comparison, 6LoWPAN is an adaptive layer defined in RFC 4944 and allows IPv6 packets to be transmitted using the short addresses in 802.15.4 (Hui and Culler, 2008). The average range for ZigBee and 6LoWPAN is 291 and 191 meters respectively (Labs, 2015).

Board		Processor			Memory		Input/Output			Voltage	Connectivity
Make	Model	Bits	Cores	Speed	App	RAM	Analogue	Digital	PWM		
Arduino	Uno	8	1	16MHz	32KB	2KB	6	14	6	5V	-
	Yún	8	1	16MHz	32KB	2.5KB	12	20	7	5V	Ethernet Wi-Fi
	Pro Mini	8	1	8MHz/16MHz	32KB	2KB	6	14	6	3.3V/5V	-
Raspberry Pi	2	32	4	900MHz	SD Card	1GB	0	40	0	5V	Ethernet
	A+	32	1	700MHz	SD Card	256MB	0	40	0	5V	-
	Zero	32	1	1GHz	SD Card	512MB	0	40	0	5V	-
Intel	Edison	64	2	500MHz	4GB	1GB	6	20	4	1.8V/3.3V	Wi-Fi a/b/g/n Bluetooth 4.0
Samsung	Artik 1	32	1	240MHz	4MB	1MB	2	0	0	Unknown	Bluetooth 4.0
	Artik 5	32	2	1GHz	4GB	512MB	2	47	2	1.8V/2.4V	Wi-Fi a/b/g/n Bluetooth 4.0 ZigBee/802.15.4
	Artik 10	32	4+4	1.3GHz/1GHz	16GB	2GB	6	51	2	1.8V/2.4V	Wi-Fi a/b/g/n Bluetooth 4.0 ZigBee/802.15.4

Table 2.3: Summary of hardware boards

2.6 Case Study: UK Smart Meters

2.6.1 Overview

The smart meter project by the British government serves as an ideal case study to exemplify an Internet of Things application. In Section 2.3, this paper described some stand-out use cases, which were collated from various sources. Upon researching the smart meter concept further, this massive governmental project—currently being developed for the United Kingdom utilities sector—was discovered. As more research on this project was carried out, it became clear that it represents a significant commentary in regard to the IoT concept.

In comparison to other potential case studies, the smart meter project is important for many reasons. First of all, the transparent and objective nature of public projects does prevent an operational bias. It would be possible to evaluate case studies provided by commercial organisations; for instance, Digi provide case studies across a number of industries. This however introduces the risk of bias—it is in the business’s interest to make a sale. In order to satisfy stakeholders, it is also necessary for the public project to unambiguously document discussions, actions and expectations. Furthermore, the scale, purpose and complexity of this project truly puts the Internet of Things concept to the test.

Before analysing the technical details of this project, it is first important to understand the purpose of smart meters. The project was announced by the Department of Energy and Climate Change (Decc) in 2009. It ambitiously aims to replace the electricity and gas meters for all residents of the United Kingdom with so-called ‘Smart Meters’ by 2020, giving a projected time frame of 10 years (The Telegraph, 2009). For consumers, these systems comprise of two main components: the electricity/gas meter and the in-home display. The display will wirelessly connect to the meter and will provide a near real-time indication for the resources being consumed, expressed in pounds and pence (The Department of Energy and Climate Change, 2009). Utility companies will also be provided with remote access to the meter, primarily allowing them take readings.

With an estimated cost of £10.9 billion (The Department of Energy and Climate Change, 2014a), the project must deliver benefits to consumers, suppliers and the government. One of the main drivers for meters are the potential financial savings. For the consumer, they are more aware of their own energy usage which allows them to reduce their spending. Into the future, this could become automated with Smart Appliances which wait for cheap energy. For the supplier, large savings will be made through the automation of meter readings and billing. Operational maintenance will also become streamlined as smart meters can immediately identify faults. For the government, efficiency savings will reduce the national environmental impact and will help the United Kingdom meet climate change targets (Thomas and Jenkins, 2012). The total projected savings from the combination of these points is £17.1 billion, resulting in a net benefit to the British economy of at least £6 billion (The Department of Energy and Climate Change, 2014a).

2.6.2 Governance

Effective governance and organisation is imperative for the success of the smart meter roll-out in the United Kingdom. Governance was discussed, in board terms, as part of Section 2.2 and found three main areas which may require governing: device addressability, information security and ethical & legal accountability. Similar issues face the British government. The Department for Energy and Climate Change have been responsible for managing the smart meter programme since April 2012 and delegates governing roles to various actors in the project. Some of the main actors include the Office of Gas and Electricity Markets (Ofgem), the purpose-built Data and

Communications Company (DCC) as well as various subcontractors and licensees. These bodies govern the smart meter project roll-out across many areas, including two which are of special interest to this paper: consumer protection and technical implementation.

Widespread social acceptance is one of the main challenges facing IoT products and services; consumer protection offered by Decc and its subcontractors goes some way to earn this public trust. First and foremost, the energy data generated by smart meters is useful to consumers, suppliers and the government. In response to the privacy issues which this creates, the Decc has developed the Data Access and Privacy Framework (DAPF). The Department of Energy and Climate Change (2015) notes that this governing framework serves three main purposes:

- Protect consumers' interests, including by addressing concerns that consumers may have about privacy;
- Enable proportionate access to data by authorised parties to ensure that benefits can be delivered; and
- Promote competition and innovation in the developing energy services market.

The DAPF is expected to be finalised by 2018. The Decc also addresses specific consumer concerns with governance and legislation. For instance, there are restrictions in regards to the supplier-led installation as detailed by the Smart Meter Installation Code of Practice, including: prohibition of sales during installation visits and the mandatory provision of energy advice.

Governance also extends to the technical standards and specifications used for the smart meter project. One important challenge facing this project is hardware and software interoperability between the large number of parties. The Smart Metering Equipment Technical Specifications (SMETS) document was drawn up in 2012 and represents the technical standards to be followed by suppliers and systems developers for the in-home products. The responsibility of Wide Area Network communication has been delegated to the Data and Communications Company (DCC). This organisation specifies its own standards for connecting to the WAN, which go as far as XML schema for inter-system communication.

2.6.3 Implementation

From the perspective of this paper, the technical implementation is the most interesting aspect of the smart meter project. There is a huge level of complexity and interconnectivity covered by SMETS and DCC specifications. The completed system will be a distributed and connected platform for smart meters in the United Kingdom and can be split into three constituent components: energy consumers, wide-area connectivity and service users. The implementation is summarised in Figure x.x.

Before discussing the implementation of specific technologies it is first worthwhile to understand the high-level system architecture. The energy consumers for the smart meter project are homes and small businesses. As was previously noted, there are three product types which each premises will have: smart meters (gas and electric), in-house displays and communication hubs. The communication hubs have two responsibilities, namely to provide premises-wide wireless connectivity—dubbed the Smart Meter Home Area Network (SM HAN)— and to connect each SM HAN to the nationwide Wide Area Network (WAN) provided by the Data and Communications Company. Service users are the utility companies like British Gas or Scottish and Southern Energy (SSE) as well as authorised third-parties such as online energy comparison websites. Authorised service users will gain bidirectional connectivity with their customer's smart meters through the DCC WAN.

The SMETS and Communications Hub Technical Specification (CHTS) documents complement one another to unambiguously define the functional requirements for smart meters and the home area networks. Since the communication hub will provide wireless connectivity across a single premises, its functional requirements are indicative of how this network will operate. The Department of Energy and Climate Change (2014b) state that the HAN interface of the communications hub will use ZigBee over the 2.4GHz frequency band. The CHTS also notes the minimum capacity of the SM HAN: four Electricity Smart Metering Equipment (ESME), one Gas Smart metering Equipment (GSME), one Gas Proxy Function (GPF), seven Type 1 devices and three Type 2 devices. Type 1 devices are those which can issue any command to ESME or GSME meters whereas Type 2 devices can only issue read requests (The Energy & Utilities Alliance, 2014).

Individual metering devices have their own functional requirements which will inform their hardware implementation. The high-level expectations of a ESME or GSME device are: a clock, a data store, meter containing one measure element, a HAN interface, a load switch, a Random Number Generator and a user interface (The Department of Energy and Climate Change, 2014c). It is also important to note that ESME or GSME devices must be capable of storing 13 months of readings taken at 30-minute intervals, even after power loss. Thomas and Jenkins (2012) compared the UK requirements against an existing smart meter, the Elster REX2. This meter has an 8-bit Texas Instrument SoC with a clock speed of 32MHz, 4KB of RAM, 256KB of application memory along with 21 GPIO lines and up to 8 analogue inputs. These specifications would place the required hardware in the realm of the Arduinos researched in 2.5 however the authors argue that the UK devices would need to be more powerful and with greater memory due to the data storage and potential computational requirements.

The final piece to the smart meter project is the Wide Area Network managed by the Data and Communications Company. The DCC is responsible for secure communication, access control and scheduled data retrieval for this project. Since the installation environments will differ from premises to premises, various connectivity methods are being employed, such as cellular, broadband and 802.15.4 mesh radio. This is, in essence, the platform supporting the whole project. The DCC will maintain a web-based service where service users can issue HTTP requests with XML document payloads to achieve tasks. The possible functions are defined in the DCC User Interface Specifications (DUIS).

2.7 Conclusions

Chapter 3

Specification

Chapter 4

Design

Chapter 5

Implementation

Chapter 6

Results and Evaluation

Chapter 7

Future Work

Chapter 8

Conclusions

Bibliography

- Divyakant Agrawal, Sudipto Das, and Amr El Abbadi. Big data and cloud computing: Current state and future opportunities. In *Proceedings of the 14th International Conference on Extending Database Technology*, EDBT/ICDT '11, pages 530–533, New York, NY, USA, 2011. ACM. ISBN 978-1-4503-0528-0. doi: 10.1145/1951365.1951432. URL <http://doi.acm.org/10.1145/1951365.1951432>.
- S. Aust, R.V. Prasad, and I.G.M.M. Niemegeers. Ieee 802.11ah: Advantages in standards and further challenges for sub 1 ghz wi-fi. In *Communications (ICC), 2012 IEEE International Conference on*, pages 6885–6889, June 2012. doi: 10.1109/ICC.2012.6364903.
- Andrew Banks and Rahul Gupta. Mqtt version 3.1.1. OASIS Standard, Oct 2014. URL <http://docs.oasis-open.org/mqtt/mqtt/v3.1.1/os/mqtt-v3.1.1-os.pdf>.
- Flavio Bonomi, Rodolfo Milito, Jiang Zhu, and Sateesh Addepalli. Fog computing and its role in the internet of things. In *Proceedings of the First Edition of the MCC Workshop on Mobile Cloud Computing*, MCC '12, pages 13–16, New York, NY, USA, 2012. ACM. ISBN 978-1-4503-1519-7. doi: 10.1145/2342509.2342513. URL <http://doi.acm.org/10.1145/2342509.2342513>.
- Case C-362/14. Maximillian Schrems v Data Protection Commissioner. OJ C351, Oct 2015. URL <http://curia.europa.eu/juris/document/document.jsf?text=&docid=172254&pageIndex=0&doclang=EN&mode=req&dir=&occ=first&part=1&cid=199525>.
- Kuor-Hsin Chang. Bluetooth: a viable solution for iot? [industry perspectives]. *Wireless Communications, IEEE*, 21(6):6–7, December 2014. ISSN 1536-1284. doi: 10.1109/MWC.2014.7000963.
- Robert Faludi. *Building Wireless Sensor Networks: With ZigBee, XBee, Arduino, and Processing*. O'Reilly Media, Inc., 1st edition, 2010. ISBN 0596807732, 9780596807733.
- I. Fette and A. Melnikov. The WebSocket Protocol. RFC 6455 (Proposed Standard), December 2011. URL <http://www.ietf.org/rfc/rfc6455.txt>.
- R. Fielding, J. Gettys, J. Mogul, H. Frystyk, L. Masinter, P. Leach, and T. Berners-Lee. Hypertext Transfer Protocol – HTTP/1.1. RFC 2616 (Draft Standard), June 1999. URL <http://www.ietf.org/rfc/rfc2616.txt>. Obsoleted by RFCs 7230, 7231, 7232, 7233, 7234, 7235, updated by RFCs 2817, 5785, 6266, 6585.
- Roy Thomas Fielding. Architectural styles and the design of network-based software architectures. Irvine, CA: University of California, 2000. URL https://www.ics.uci.edu/~fielding/pubs/dissertation/rest_arch_style.htm.

- Github, Inc. Choose a license. San Francisco, CA: Github, Oct 2015. URL <http://choosealicense.com/>. [Accessed 14 October 2015].
- Jonathan W. Hui and David E. Culler. Ip is dead, long live ip for wireless sensor networks. In *Proceedings of the 6th ACM Conference on Embedded Network Sensor Systems*, SenSys '08, pages 15–28, New York, NY, USA, 2008. ACM. ISBN 978-1-59593-990-6. doi: 10.1145/1460412.1460415. URL <http://doi.acm.org/10.1145/1460412.1460415>.
- Dogan Ibrahim. Advanced pic microcontroller projects in c. Oxford: Newnes, 2011. URL <http://www.myilibrary.com?ID=127324>. [Accessed 1 December 2015].
- Link Labs. M2m & iot technologies explained. Annapolis, MD: Link Labs, 2015.
- Kalle Lyytinen and Youngjin Yoo. Introduction. *Commun. ACM*, 45(12):62–65, December 2002. ISSN 0001-0782. doi: 10.1145/585597.585616. URL <http://doi.acm.org/10.1145/585597.585616>.
- L. Masinter. Hyper Text Coffee Pot Control Protocol (HTCPCP/1.0). RFC 2324 (Informational), April 1998. URL <http://www.ietf.org/rfc/rfc2324.txt>. Updated by RFC 7168.
- Friedemann Mattern and Christian Floerkemeier. From active data management to event-based systems and more. chapter From the Internet of Computers to the Internet of Things, pages 242–259. Springer-Verlag, Berlin, Heidelberg, 2010. ISBN 3-642-17225-3, 978-3-642-17225-0.
- K. Mikhaylov and J. Tervonen. Optimization of microcontroller hardware parameters for wireless sensor network node power consumption and lifetime improvement. In *Ultra Modern Telecommunications and Control Systems and Workshops (ICUMT), 2010 International Congress on*, pages 1150–1156, Oct 2010. doi: 10.1109/ICUMT.2010.5676525.
- J. Mineraud, O. Mazhelis, X. Su, and S. Tarkoma. Contemporary Internet of Things platforms. *ArXiv e-prints*, January 2015.
- Julien Mineraud and Sasu Tarkoma. Toward interoperability for the internet of things with meta-hubs. Technical Report arXiv:1511.08063v1, University of Helsinki, Helsinki, Finland, nov 2015. URL <http://arxiv.org/pdf/1511.08063v1.pdf>.
- Luca Mottola and Gian Pietro Picco. Programming wireless sensor networks: Fundamental concepts and state of the art. *ACM Comput. Surv.*, 43(3):19:1–19:51, April 2011. ISSN 0360-0300. doi: 10.1145/1922649.1922656. URL <http://doi.acm.org/10.1145/1922649.1922656>.
- J. Postel. User Datagram Protocol. RFC 768 (INTERNET STANDARD), August 1980. URL <http://www.ietf.org/rfc/rfc768.txt>.
- Z. Shelby, K. Hartke, and C. Bormann. The Constrained Application Protocol (CoAP). RFC 7252 (Proposed Standard), June 2014. URL <http://www.ietf.org/rfc/rfc7252.txt>.
- P. Srisuresh and M. Holdrege. IP Network Address Translator (NAT) Terminology and Considerations. RFC 2663 (Informational), August 1999. URL <http://www.ietf.org/rfc/rfc2663.txt>.
- The Department for Business Innovation & Skills. Smart cities background paper. London: Department for Business Innovation & Skills Publications, Oct 2013. URL https://www.gov.uk/government/uploads/system/uploads/attachment_data/file/246019/bis-13-1209-smart-cities-background-paper-digital.pdf.

- The Department of Energy & Climate Change. Smart meters: a guide. *London: Department of Energy & Climate Change*, Oct 2013. URL <https://www.gov.uk/guidance/smart-meters-how-they-work>. [Accessed 29 November 2015].
- The Department of Energy and Climate Change. Smart meters: a guide. London: Department of Energy and Climate Change [Online], Oct 2009. URL <https://www.gov.uk/guidance/smart-meters-how-they-work>.
- The Department of Energy and Climate Change. Third annual report on the roll-out of smart meters. London: Department of Energy and Climate Change, Dec 2014a.
- The Department of Energy and Climate Change. Communications hub technical specifications. London: Department of Energy and Climate Change, Nov 2014b.
- The Department of Energy and Climate Change. Smart metering equipment technical specifications. London: Department of Energy and Climate Change, Nov 2014c.
- The Department of Energy and Climate Change. Consultation on the timing of the review of the data access and privacy framework. London: Department of Energy and Climate Change [Online], Mar 2015. URL <https://www.gov.uk/government/consultations/consultation-on-the-timing-of-the-review-of-the-data-access-and-privacy-framework>.
- The Energy & Utilities Alliance. Smart metering device assurance scheme operator services. Kenilworth: Energy & Utilities Alliance [Online], Sep 2014. URL [http://www.eua.org.uk/sites/default/files/SMDA%20Scheme%20Operator%20RFP%20v1.2%20\(Final\).pdf](http://www.eua.org.uk/sites/default/files/SMDA%20Scheme%20Operator%20RFP%20v1.2%20(Final).pdf).
- The European Commission. Freight transport logistics action plan. *Luxembourg: Publications Office of the European Union, COM/2007/0607*, Oct 2007. URL <http://www.ipex.eu/IPEXL-WEB/dossier/dossier.do?code=COM&year=2007&number=0607>.
- The European Commission. Action plan for the deployment of internet protocol version 6 (ipv6) in europe. *Luxembourg: Publications Office of the European Union, COM/2008/0313*, May 2008. URL <http://www.ipex.eu/IPEXL-WEB/dossier/document/COM20080313FIN.do>.
- The Telegraph. Energy suppliers to install smart meters in all households by 2020. Dec 2009. URL <http://www.telegraph.co.uk/news/earth/energy/6708822/Energy-suppliers-to-install-smart-meters-in-all-households-by-2020.html>. [Accessed 5 December 2015].
- Lee Thomas and Prof. Nick Jenkins. Smart metering for the uk. HubNet [Online], Jun 2012. URL http://www.hubnet.org.uk/position_papers.
- J. Wei. How wearables intersect with the cloud and the internet of things : Considerations for the developers of wearables. *Consumer Electronics Magazine, IEEE*, 3(3):53–56, July 2014. ISSN 2162-2248. doi: 10.1109/MCE.2014.2317895.
- Mark Weiser. The computer for the 21st century. *SIGMOBILE Mob. Comput. Commun. Rev.*, 3(3):3–11, July 1999. ISSN 1559-1662. doi: 10.1145/329124.329126. URL <http://doi.acm.org/10.1145/329124.329126>.